

NovaCore AI Systems - Full Product, Billing, Technical & Security Documentation

1. Product Description

NovaCore AI Systems is a cloud-native AI platform designed to help enterprises build, deploy, and manage intelligent agents. The system integrates LLMs, APIs, memory, and workflow orchestration into a unified environment.

1.1 Core Product: NovaCore Agent Platform

The Agent Platform allows users to create AI agents using a modular system combining prompts, tools, and workflows. It supports multi-model inference, including OpenAI-compatible APIs and open-source LLMs. Agents can perform tasks such as automation, reasoning, customer support, and data extraction.

Key capabilities include drag-and-drop workflow design, tool chaining, memory persistence using vector databases, and integration with external APIs such as CRMs, databases, and cloud services.

1.2 NovaCore Monitor

NovaCore Monitor provides observability for AI systems. It tracks latency, token usage, cost, and system behavior. It includes detection of unsafe outputs, hallucinations, and prompt injection attempts.

The system generates dashboards for engineering teams and alerts when abnormal behavior is detected. It is designed to ensure production-grade safety and reliability.

1.3 Billing Intelligence Module

Billing Intelligence tracks real-time token usage and API consumption across all agents. It provides forecasting, budget alerts, and optimization recommendations to reduce operational costs.

2. Billing & Subscription Policy

NovaCore uses a hybrid subscription and usage-based billing model. Customers are charged based on compute usage, token consumption, and active features.

2.1 Pricing Tiers

Free Tier includes limited tokens, one agent, and basic monitoring. Pro Tier provides expanded token limits, up to 10 agents, and advanced analytics. Enterprise Tier includes unlimited usage, SLA guarantees, and dedicated infrastructure.

2.2 Billing Cycle

Billing is calculated daily based on usage metrics and invoiced monthly. Token usage, API calls, and tool execution costs are aggregated into a unified bill.

2.3 Refund Policy

Subscriptions may be refunded within 7 days of initial purchase. Usage-based charges are non-refundable once processed. Enterprise contracts follow negotiated agreements.

2.4 Payment Processing

Payments are processed via secure third-party providers. Supported methods include credit cards and ACH transfers for enterprise customers.

3. Technical Troubleshooting Guide

This section describes common system issues and recommended debugging procedures.

3.1 Agent Failures

Agent failures often result from invalid API keys, rate limiting, or model unavailability. Engineers should verify authentication, inspect logs, and check system health dashboards.

3.2 Latency Issues

High latency is typically caused by large context windows, slow external tool calls, or overloaded inference endpoints. Optimization includes prompt reduction and caching strategies.

3.3 Billing Discrepancies

Billing mismatches are usually due to delayed telemetry sync or overlapping agent logs. Reconciliation jobs should be executed to validate usage records.

3.4 RAG Failures

Retrieval-Augmented Generation failures occur when vector databases are not indexed or embeddings are mismatched. Re-indexing and model alignment are required.

4. Security & Policy Documentation

NovaCore enforces strict security policies to ensure safe deployment of AI systems in production environments.

4.1 Data Security

All data is encrypted at rest using AES-256 and in transit using TLS 1.2+. Customer data is not used for training unless explicitly authorized.

4.2 Prompt Injection Defense

The system includes multi-layer defenses against prompt injection including input sanitization, system prompt isolation, and output filtering.

4.3 Access Control

Role-based access control ensures least privilege access. Administrative actions require multi-factor authentication and are logged for audit purposes.

4.4 Incident Response

Incidents are classified by severity from SEV1 to SEV4. SEV1 incidents require immediate response within 15 minutes and escalation to engineering leadership.

4.5 Acceptable Use

Users are prohibited from attempting jailbreaks, extracting system prompts, bypassing safety layers, or abusing system resources. Violations may result in suspension or termination.